

# Practical Machine Learning

## Day 6: Sep22 DBDA

Kiran Waghmare

# Agenda

- Ridge, Lasso & ElasticNet
- Preprocessing Techniques
- Classification

File

Edit

View

Insert

Cell

Kernel

Widgets

Help

Not Trusted

Python 3 (ipykernel)



Code



Commit



```
In [37]: from sklearn.linear_model import LinearRegression
reg=LinearRegression()
reg.fit(X_train, y_train)
```

```
Out[37]: LinearRegression()
```

## Prediction

```
In [38]: y_pred=reg.predict(X_test)
y_pred
np.set_printoptions(precision=2)
print(np.concatenate((y_pred.reshape(len(y_pred),1), y_test.reshape(len(y_test),1)),1))
```

```
[[103015.2 103282.38]
 [132582.28 144259.4 ]
 [132447.74 146121.95]
 [ 71976.1   77798.83]
 [178537.48 191050.39]
 [116161.24 105008.31]
 [ 67851.69  81229.06]
 [ 98791.73  97483.56]
 [113969.44 110352.25]]
```

Simple  
Linear  
Regression


$$y = b_0 + b_1 x_1$$

Multiple  
Linear  
Regression


$$y = b_0 + b_1 x_1 + b_2 x_2 + \dots + b_n x_n$$

Polynomial  
Linear  
Regression


$$y = b_0 + b_1 x_1 + b_2 x_1^2 + \dots + b_n x_1^n$$



**Underfitting**




**Just right!**




**overfitting**


Linear





Linear



No linear relationship



Copyright 2014. Laerd Statistics.

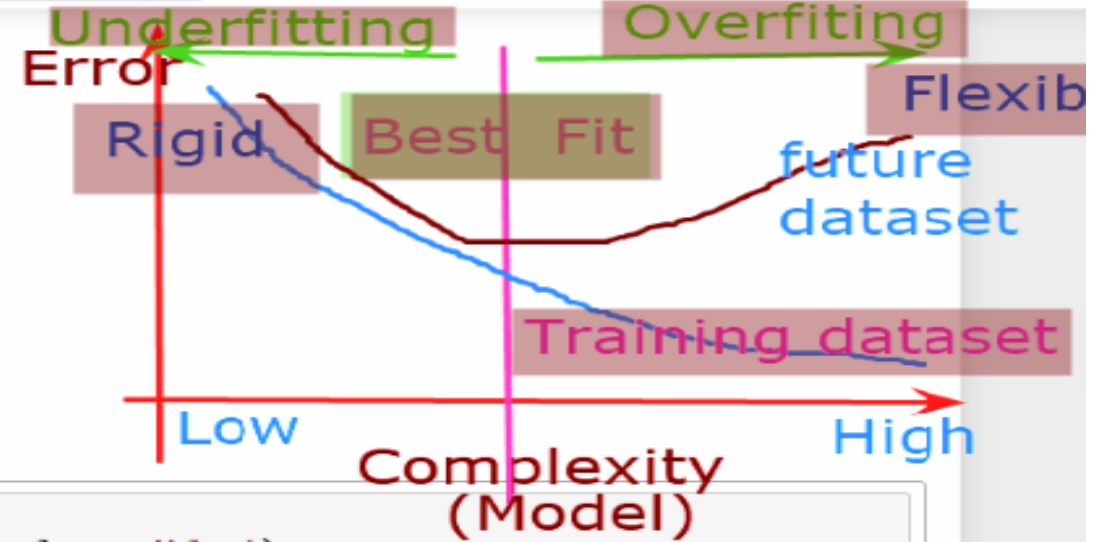





A photograph of a wooden bed frame on a light-colored wooden floor. Instead of a standard mattress, a mattress is placed on the frame that is shaped like the number '4'. The mattress is white with a quilted pattern. The bed frame is dark wood with decorative posts at the corners. The background is a plain white wall.



**THE BEST WAY TO  
EXPLAIN OVERFITTING**



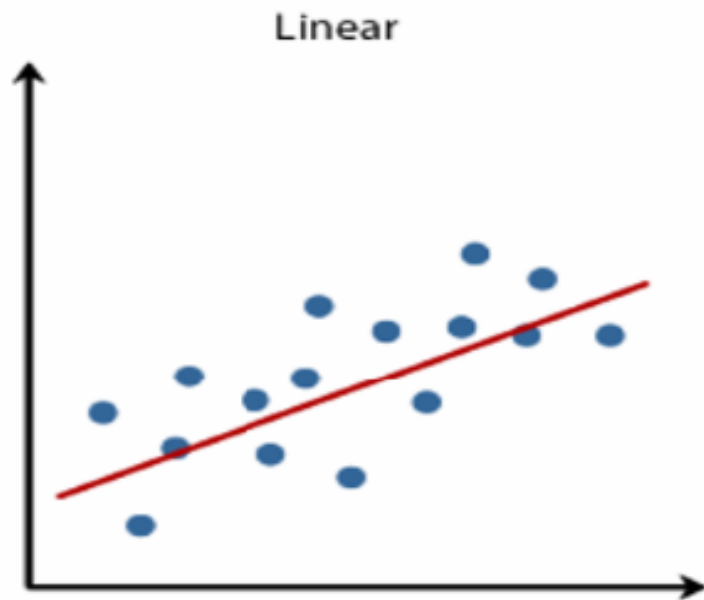




```
In [39]: plt.scatter(X,y, color = 'red')
plt.plot(X,model1.predict(poly_reg.fit_transform(X)), color = 'blue')
```

Out[39]: [<matplotlib.lines.Line2D at 0x2b561d8e160>]



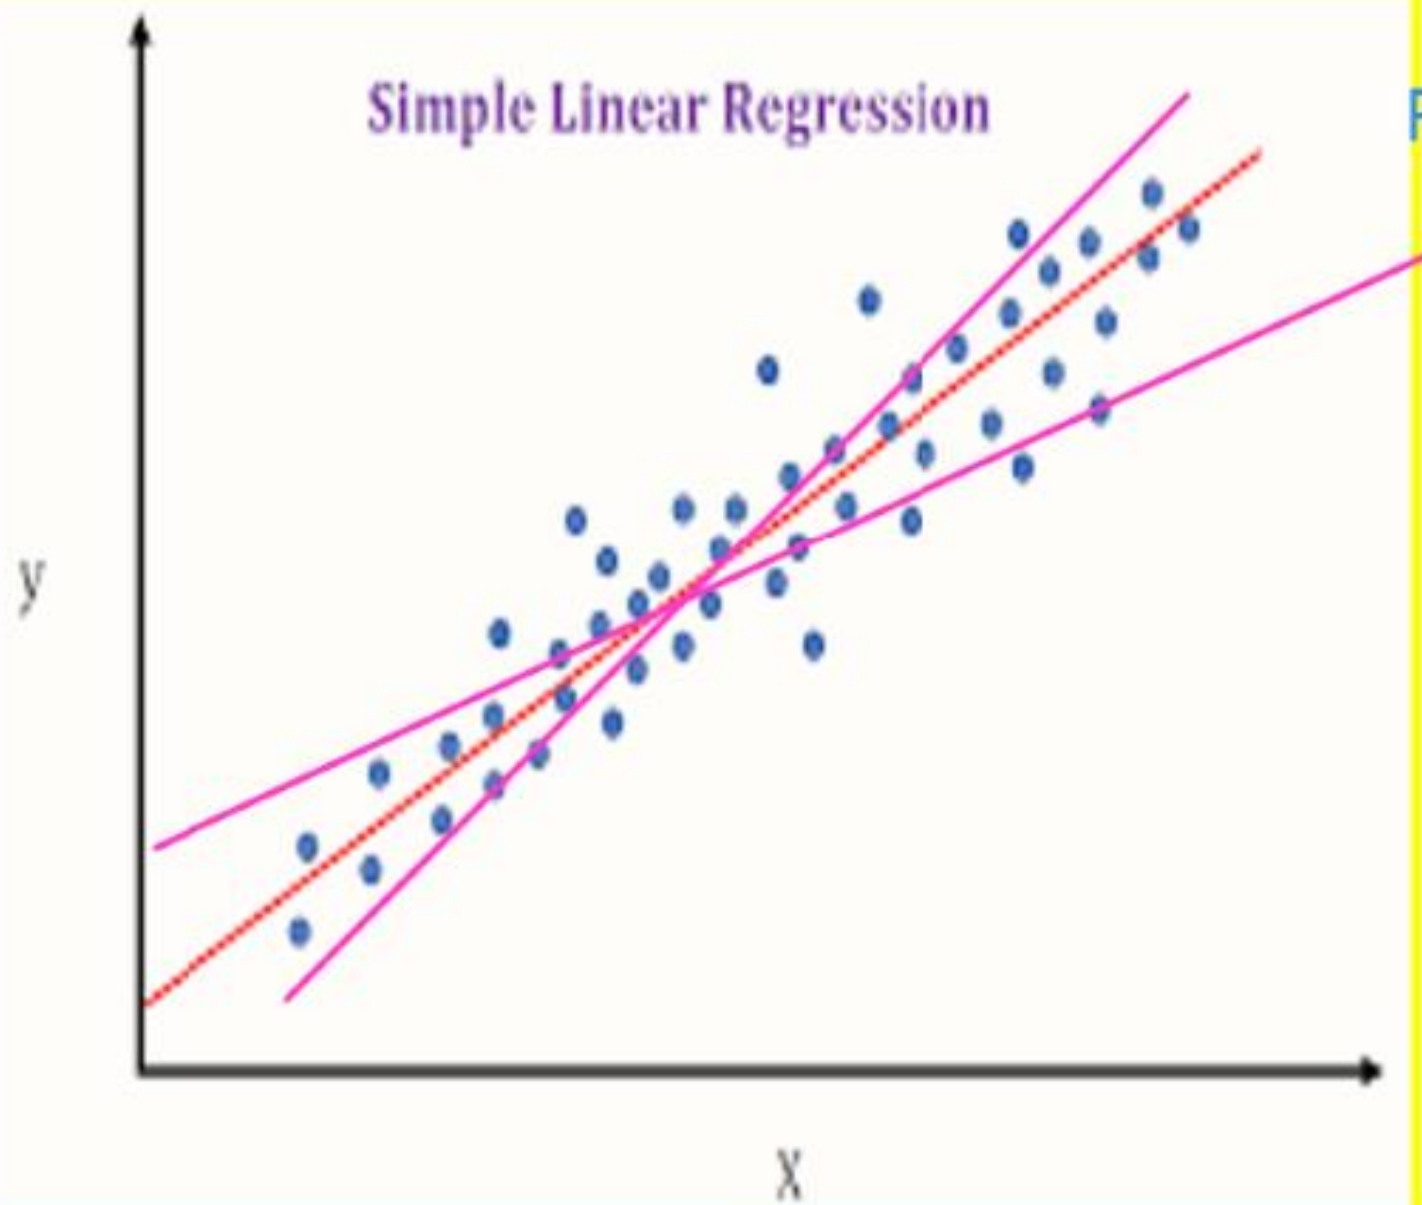


Copyright 2014. Laerd Statistics.

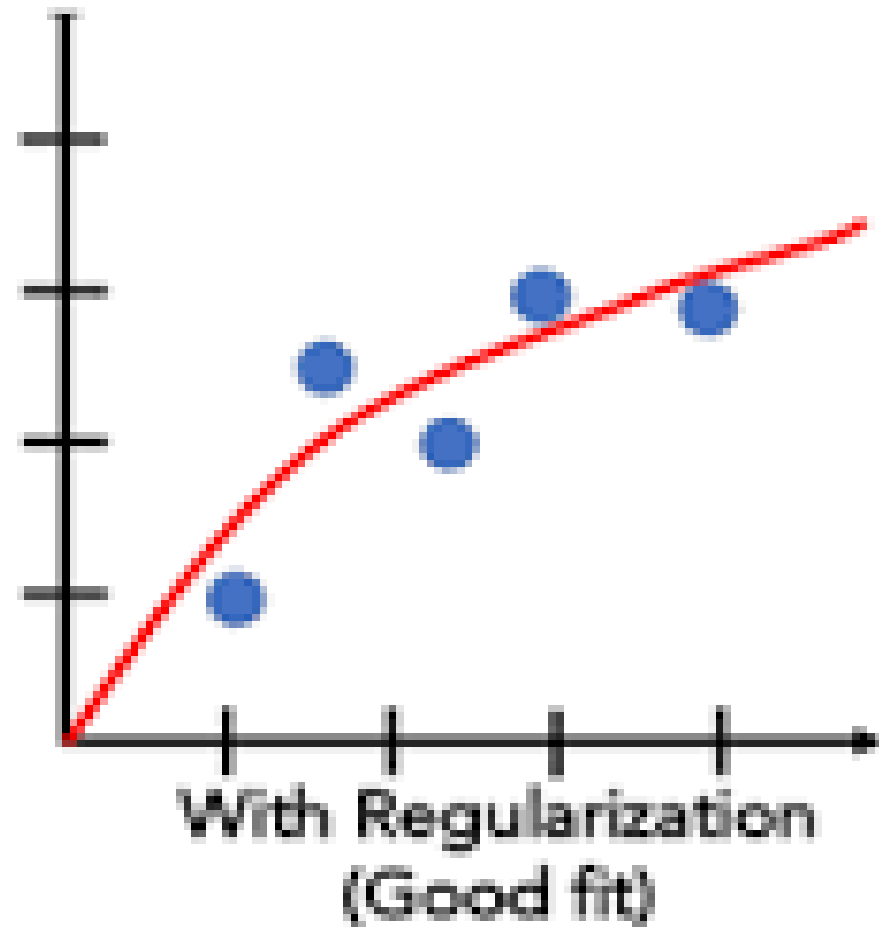

Overfitting ---> Good fit  
(Regularization)

Simple Linear Regression

Regression line(best fit)



## *Impact of Regularization*



# Regularization

**Regularization Term**

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m \left( h_{\theta}(x^{(i)}) - y^{(i)} \right)^2 + \boxed{\lambda \sum_{j=1}^n \theta_j^2}$$

Regularization Parameter

start at  $\theta_1$

# Regularization

x1	x2	x3	x4	y
		✓	✓	↓

Penalty

Regularization Term

$$J(\theta) = \frac{1}{2n} \sum_{i=1}^m \left( h_{\theta}(x^{(i)}) - y^{(i)} \right)^2 + \lambda \sum_{j=1}^n \theta_j^2$$

Regularization Parameter


start at  $\theta_1$

x1:1  
x2:1.5  
x3:90  
x4:85

Ridge  
Lasso  
ElasticNet

Loss function + Regularized term





```
In [12]: new_data['Houseprice'] = data.target
new_data.head()
```

```
Out[12]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	Houseprice	House_price
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98	24.0	24.0
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14	21.6	21.6
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03	34.7	34.7
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94	33.4	33.4
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33	36.2	36.2

```
In [ ]:
```

```
4 0.06905 0.0 2.18 0.0 0.458 7.147 54.2 6.0622 3.0 222.0 18.7 396.90 5.33 36.2
```

```
In [23]: new_data.shape
```

```
Out[23]: (506, 14)
```

```
In [24]: x=new_data.iloc[:,0:13]  
y=new_data.iloc[:,13]
```

```
In [25]: x
```

```
Out[25]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33
...	...	...	...	...	...	...	...	...	...	...	...	...	...
501	0.06263	0.0	11.93	0.0	0.573	6.593	69.1	2.4786	1.0	273.0	21.0	391.99	9.67
502	0.04527	0.0	11.93	0.0	0.573	6.120	76.7	2.2875	1.0	273.0	21.0	396.90	9.08
503	0.06076	0.0	11.93	0.0	0.573	6.976	91.0	2.1675	1.0	273.0	21.0	396.90	5.64
504	0.10959	0.0	11.93	0.0	0.573	6.794	89.3	2.3889	1.0	273.0	21.0	393.45	6.48
505	0.04741	0.0	11.93	0.0	0.573	6.030	80.8	2.5050	1.0	273.0	21.0	396.90	7.88

506 rows × 13 columns

# Regularization

Regression:

- Ridge
- Lasso

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m \left( h_{\theta}(x^{(i)}) - y^{(i)} \right)^2 + \underbrace{\lambda \sum_{j=1}^n \theta_j^2}_{\text{Regularization Term}}$$

Regularization Parameter

start at  $\theta_1$

Loss

# Ridge Regression

$$\text{Regularization (L2)} = \text{Loss Function} + \lambda \sum_{i=1}^m w_i^2$$


# Ridge Regression

$$\text{Regularization (L2)} = \text{Loss Function} + \lambda \sum_{i=1}^m w_i^2$$

# Lasso Regression

$$\text{Lasso Regression} = \operatorname{argmin}_{\beta \in \mathbb{R}^p} \underbrace{\|y - X\beta\|_2^2}_{\text{Loss}} + \lambda \underbrace{\|\beta\|_2^2}_{\text{Penalty}}$$





# Regularization

$$y = 0.8 + 0.6x_1 + 0.4x_2 + 1232434x_3$$

cost  
time  
computation

Regularization Term

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m \left( h_{\theta}(x^{(i)}) - y^{(i)} \right)^2 + \lambda \sum_{j=1}^n \theta_j^2$$

start at  $\theta_1$

Regularization Parameter

$\lambda = 0.01$


$\lambda = 1$

$\lambda = 2$

```
Out[21]: ridge()
```


```
In [22]: rid_coeff=pd.Series(rid.coef_,index=data.feature_names)  
rid_coeff
```

```
Out[22]: CRIM      -0.962257  
ZN          1.040872  
INDUS       0.011680  
CHAS        0.598719  
NOX         -1.820134  
RM          2.583786  
AGE         -0.095188  
DIS         -2.848263  
RAD          2.036231  
TAX         -1.806092  
PTRATIO     -2.283191  
B           0.718310  
LSTAT       -3.576073  
dtype: float64
```



```
In [23]: rid_coeff.plot(kind='bar')
```

```
Out[23]: <AxesSubplot:>
```



```
In [ ]: #Lasso Algorithm  
from sklearn.linear_model import Lasso
```

```
In [25]: la_coff = pd.Series(la.coef_, index=data.feature_names)
la_coff
```

```
Out[25]: CRIM    -0.000000
          ZN      0.000000
          INDUS  -0.000000
          CHAS    0.000000
          NOX    -0.000000
          RM      2.540098
          AGE    -0.000000
          DIS    -0.000000
          RAD    -0.000000
          TAX    -0.171527
          PTRATIO -1.784796
          B       0.110959
          LSTAT   -3.585324
          dtype: float64
```

Feature Selection

Regularization

Lasso

+

Ridge

Elasticnet

Multi feature correlation

```
In [26]: la_coff.plot(kind='bar')
```

```
Out[26]: <AxesSubplot:>
```




# Transforming the Loss function into Lasso Regression

$$\sum_{i=1}^n \left( y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 \Rightarrow \sum_{i=1}^M \left( y_i - \sum_{j=0}^p w_j \times x_{ij} \right)^2 + \lambda \sum_{j=0}^p |w_j|$$

Loss function

Loss function + Regularized term

Designed by Author (Shanthababu)



# Ridge Regression

Ridge regression uses the mean squared error loss function and applies L2 Regularization. Its cost function  $J(\theta)$  is given as

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m (y - \hat{y})^2 + \lambda \sum_{j=1}^n w_j^2$$

where,

$\frac{1}{m} \sum_{i=1}^m (y - \hat{y})^2$  is the Mean Squared error (loss function)

$\lambda \sum_{j=1}^n w_j^2$  is the penalty (L2 Regularization)

Now, substitute  $\hat{y}$  as  $w x_i + b$ .



# Lasso Regression

Lasso regression uses the same mean squared error loss function and this applies L1 Regularization and will repeat the same steps as Ridge. The cost function of Lasso Regression  $J(\theta)$  is given as

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m (y - \hat{y})^2 + \lambda \sum_{j=1}^n |w_j|$$

where

$\lambda \sum_{j=1}^n |w_j|$  is the penalty (L1 Regularization).